

FOUR DAYS OF WORK

Everything designed, built, broken, diagnosed and fixed
between 27 and 30 August 2026 — counted from the
commit log, not from memory.

185

COMMITTS

154

FILES TOUCHED

+26,862


LINES ADDED

-1,711

LINES REMOVED

AT A GLANCE


Four days, 185 commits, 154 files — 113 of them pair-built with Claude. The shape of the work: roughly one part new capability to one part correction, on a base of continuous housekeeping.



COMMITTS PER DAY



COMMITTS BY TYPE



HOW TO READ THE CHORE COUNT

Chore is the largest single bucket, and that is expected rather than alarming: it covers release tagging, dependency bumps, catalogue refreshes and the automated daily digests. It is the cost of keeping a live site running, not work done instead of features.

EVERYTHING HERE IS ON THE LIVE SITE

Every commit counted in this report is merged to **master** and deployed. `toolnaut.xyz` returns HTTP 200 on the current build, and all 20 public routes render with zero console errors.

WHAT SHIPPED

32 features, in the words of the commits that introduced them.

- [galaxy] daily radar arrivals rise as NAMED stars, automatically
 - [claims] claim_key normaliser with collision detection, ahead of first import
 - [footer] full-width mark on a charcoal panel — the user’s sketch, literally
 - [footer] the mark is now the BACKGROUND of the contact section; hero gets air
 - [footer] the mark idles blurred in the middle; the cursor pulls it sharp
 - [brand] one lockup everywhere — beta tag built in, sized up, hero deduped
 - [cursor] Plasma Ribbon by default; all ten effects + size in ME settings
 - [api] harden /api/chat; prep the tool_claims provenance migration
 - [goal] open with the badge from direction A, not the wordmark (#14)
 - [scripts] cursor-effect picker — 10 candidate effects + demo assembler
- [sync] server-backed persistence, inert until the schema is applied
 - [seo] per-route titles and canonicals — the prerender was being thrown away
 - [landing] count explorers for real, and show nothing when it cannot
 - [build] prerender the public routes to real HTML
 - [app] open /app to guests — the sign-in gate guarded nothing
 - public /new feed surfaces radar-discovered tools without a session
 - [profile] add the 16-character explorer crew as profile avatars
 - [app] production-polish pass across all six app sections
 - [auth] port the cabinet silhouette, and make the deck buttons work
 - [auth] put Toolnaut’s real galaxy on the cabinet screen, flown by the joystick
 - [auth] port the cabinet trim from the approved sign-in build
- the infinity wordmark, and a hero reveal that draws it
 - add public category landing pages (“best AI tools for X”)
 - [auth] detail the cabinet to match the reference
 - real privacy and terms pages, written from what the app actually does
 - arcade-cabinet sign-in modal, with a screen that is actually running
 - real Google sign-in via Supabase, behind a simulated fallback
 - explorer/subscriber figures on the landing, and a real rank in the app
 - two skies — full moon and moonless — and a denser starfield
 - remove /starchart, and with it the second price list
 - landing gets live numbers, subscriptions and a leaderboard
 - let Naut actually understand a typed reply, via Featherless

THE FOUR THAT MATTERED MOST

FEATURE	WHY IT COUNTED
Open the app to guests	The sign-in gate was guarding a product that works fine without an account. Removing it turned the app itself into the demo.
Prerendered public routes	The site was shipping an empty shell to crawlers. Public routes are now real HTML with per-route titles and canonicals.
The arcade sign-in cabinet	The sign-in screen became a working cabinet: the real galaxy on the CRT, flown by the joystick, with the deck buttons wired up.
Server-backed sync	User state can move off this browser. Deliberately inert until the schema is applied — see the open items.

WHAT WAS FIXED

34 fixes. Several were bugs that no test would have caught, because the code was correct and the *output* was wrong.

- [app] guests saw a 'Plan:' label with nothing after it
- [pricing] reconcile PricingSection's unqualified claims with CapabilityMatrix's honest live/planned split
- [picker] only offer each sky control where it does something (#19)
- [goal] badge and headline move INSIDE the chat box; bigger logo; no start gap
- [goal] the headline carries the sample's weight — two single strokes
- [footer] sink the mark into the page's bottom edge; least possible area
- [footer] remove the bottom bar — the mark is now the page's last word
- [footer] one flat Obsidian #020403; compact the section; mark stays last
- [footer] the reveal now shows the actual mark, not a dotted tracing of it
- [cursor] the fade blacked out the entire app — erase alpha, don't paint dark
- [api] anchor preview-origin regexes; normalise claim_key at the database
- [api] tighten the Vercel-preview origin check; precise claim dedup identity
- [security+ux] close the open redirect, add security headers, real 404, first src/ tests
- [pricing] finish flattening the row; make smoke immune to cold-start flake
- [state] scope stored data per account, and offer to import guest work
- [build] keep three.js off the entry critical path
- let a shared stack actually be adopted, not just viewed
- [goal] make the intake a fixed chat frame the transcript scrolls inside
- [build] install chromium's system libs on Vercel before prerendering
- [build] make the prerender actually run on Vercel, and prove it
- [build] commit utils/funnel.js — master could not build without it
- [app] finish the heading pass, and make the no-persona SAVED page do something
- [discover] use the width a wide screen actually has
- [auth] drop the GitHub button — the provider is disabled, so it always failed
- [stack] rebuild the first-run screen, the app's actual front door
- add skip-to-content link to app shell
- hero badge still says "launching soon" on a live public beta
- [radar] move schedule out of the 00:00-02:00 UTC window, run twice daily
- quiz-length copy said two minutes, everywhere else says 60 seconds
- take the mascot out of the hero wordmark
- drop the landing leaderboard, and make the community row readable
- [auth] sign-in sent new accounts to an empty dashboard
- [auth] signing in with Google landed you back on the sign-in page
- [radar] move scheduled run off the top of the hour

THE CURSOR EFFECT BLACKED OUT THE ENTIRE APP

A trail effect faded by painting translucent black over the canvas each frame. On a dark UI that accumulated into an opaque sheet covering the page. The fix was to erase alpha instead of painting darkness — the same visual intent, the opposite operation.

A BUILD THAT WAS GREEN AND STILL BROKEN

A scripted edit inserted markup containing the word `BrandLogo`, and the guard that was supposed to add the missing import checked whether that word appeared anywhere in the file. It now did, so the import was skipped on three pages. The build stayed green; the pages threw at runtime. Only a render audit caught it. The guard now matches an actual import statement, not a substring.

A WRONG CONCLUSION, CORRECTED

The joystick was reported as not moving the cabinet's galaxy. That verdict came from frame-differencing a near-black CRT, which was far too insensitive to detect it. Direct instrumentation showed the camera does move — resting at $y = 4.10$ and rising to $y = 8.30$ under stick input. The measurement was at fault, not the feature.

BEFORE & AFTER

Captured from the running app on the same routes, at the same viewport, before and after the work.


PRICING

Unqualified plan claims were reconciled against what the app can actually do today; anything not yet built is now marked planned rather than sold.

BEFORE




AFTER




INTAKE (/GOAL)

The intake became a fixed chat frame that the transcript scrolls inside, with the wordmark centred in the header and the progress bar removed.

BEFORE




AFTER




YOUR STACK

The first-run screen — the app’s real front door — was rebuilt, and the whole shell got a production-polish pass.

BEFORE




AFTER




TOOL DETAIL

Detail pages gained the trust panel that explains why a tool was matched to your answers instead of asserting a score.

BEFORE



AFTER



QUALITY GATES

Every push runs the same three gates. Nothing in this report was merged without all three passing.

GATE	WHAT IT PROVES	STATUS
Unit tests	162 assertions across scoring, state stores, catalogue integrity and the radar pipeline	102 + 60 pass
Route smoke	All 20 public routes boot in a real browser and report zero console errors	20 / 20 clean
Production build	Vite build plus the prerender pass that turns public routes into real HTML	Green

STILL OPEN

Two items are finished in code and blocked on something outside the repository. Both are listed here rather than quietly counted as done.

ITEM	STATE
Supabase migrations 0001 → 0003	Written, reviewed and committed. Not yet applied to the database, so server-backed sync stays inert by design and all user data continues to live in this browser. Applying them is a console action, not a code change.
VITE_GA4_ID	Analytics events are wired throughout the app and fire correctly. The measurement id is not set in Vercel Production, so nothing is being recorded yet.

ONE RISK WORTH NAMING

The autonomous routines that keep the catalogue fresh have more than once overwritten directed work — reverting the centred wordmark on the intake screen and rewriting scene files after they had been set deliberately. Anything hand-tuned should be re-checked after a routine runs.